



Module IN3031 / INM378

Digital Signal Processing and Audio Programming

Johan Pauwels

johan.pauwels@city.ac.uk

based on slides by Tillman Weyde



(Fast) Fourier Transform in Practice



Summary: DFT

- The **Discrete Fourier Transform** calculates the **spectrum X** from a **signal x** :

$$X[k] = \sum_{n=0}^{N-1} x[n] e^{-i \frac{2\pi}{N} nk}$$

- The **inverse** is very similar

$$x[n] = \frac{1}{N} \sum_{k=0}^{N-1} X[k] e^{+i \frac{2\pi}{N} nk}$$



Summary: FFT

- The **Fast Fourier Transform** (fft,ifft) is **efficient**: $O(N \log N)$ instead of $O(N^2)$
- Used on **signals** of **length n** being **power of 2**
- Signals can be **zero-padded** to fit



FFT Normalisation

- The FFT **preserves the energy** contained in a signal, if multiplied by **factor** \sqrt{N}
- The **whole FFT / iFFT cycle** produces a **multiplication** by N
- Normally **cancelled** only **in the iFFT** for efficiency reasons
- Alternatively, **divide** the **power spectrum** by N .



Magnitude and Power Spectrum

- The **magnitude** (absolute of complex number) **spectrum** describes the **amplitude** for each frequency used (so called **bins**).
- The **square** of the **absolute** value describes the **energy** of the signal in that **bin**.



Short-Time FFT Windowing Convolution & Filtering



Spectra over Time

- Even **FFT** takes long to compute for a long signal
- In **real time** we want frequency information before the signal has ended
- Often, **very low frequencies** are **not of interest**
- Different **spectra over time** are of interest (**one FFT** gives **one spectrum**)




Short Time Fourier Transform

- **Fourier transform** performed on **time windows**, i.e. short parts of the signal
- **Sequence** of window **spectra** describes the spectral development over time
- Result: **2-dimensional** structure
time vs. frequency

Spectrogram

- The resulting 2D image is called a **spectrogram**.





Window Length

- The **length** of the **window** determines the
 - **Lowest** represented **frequency**
 - **Resolution** of the **spectrum**
- **Trade-off** between **time** and **frequency resolution** (Heisenberg's uncertainty principle)
 - **Long windows** offer more information in the spectra (better **frequency resolution**).
 - **Short windows** offer more information on the time-scale (better **time resolution**).



Problems with Windowing

- Simplest window: **rectangular**

$$w_R(n) = \begin{cases} 1 & n = 0, \dots, R-1 \\ 0 & \text{elsewhere} \end{cases}$$

- **Fourier Transform requires periodic signals.**
- **Signals are usually not periodic, or not at window length period.**








Better Windows

- A **window** function that **fades to 0 smoothly** **avoids problems** with non-periodic signals.
- Smooth window **reduces frequency resolution**.
- **Common choice**: the **Hann** window, defined as:

$$w[n] = \frac{1}{2} - \frac{1}{2} \cos(2\pi n/N)$$

Hann Window


- The Hann window is available in Python (SciPy) as `scipy.signal.windows.hann(N)`





Spectrum Leaking

- The FT of the window function shows how the **windowed spectrum 'leaks'** for different window functions (red: Hann, green: Rectangular).





Python Functions for FFT

Get the **spectrum** (complex numbers)

`spec = scipy.fft.rfft(sig) or scipy.fft.fft(sig)`

Use `np.abs(spec)` for **magnitude** and `np.angle(spec)` for **phase**.

Apply `np.real()` and `np.imag()` if needed (rare).

Get the signal back with

`sig = scipy.fft.irfft(spec) or scipy.fft.ifft(sig)`



Resynthesis from the Spectrogram



Resynthesis

- From the STFT we can **return** to the **signal** by **iFFT** on **every window**.
- ... but we **can change** it before doing that :-)
- interesting applications:
 - **equalizing**: just amplify or damp
 - **denoising**: subtract the noise-floor
 - **watermarking**: imprint a specific pattern
 - **vocoding**: transferring spectrum peaks



FFT Normalisation

- The **signal** (s) has the same **energy** as the **spectrum** (S) divided by \sqrt{N} $\sum s^2 = \sum \left(\frac{S}{\sqrt{N}} \right)^2$
(this is called *Parseval's theorem*)
- Alternatively, divide the power spectrum (i.e. S^2) by N to achieve normalisation: $\sum s^2 = \frac{1}{N} \sum S^2$
- The whole **FFT / iFFT** cycle produces a multiplication by N which is **normally cancelled only** in the **iFFT** for efficiency reasons.




Reconstruction after Windowing

- **Windowing** means $y[n] = x[n] * w[n]$
- **Reverse** after STFT by $x'[n] = y[n] / w[n]$
- **Problems:**
 - **rounding errors** near the margins
 - **large values** near the margins **if** the **spectrum was processed**
- **Idea: let windows overlap**



Getting the Joints Right

- Windows must **overlap** to get good quality.
- If the **signal** was **changed**, there will be **clicks**:






Windowing and Crossfading

- **Crossfading**: blending from one signal to another.
- Idea: use **overlapping** windows, that fade in and out in the overlap region.
- Window should be designed so that values **add up to 1** in the **overlap range**.



Crossfading Functions

- Condition: Function ranges w_{in} and w_{out} should **add up to 1** in the **overlap**.
- E.g. **triangle** waves or a **Hann** window:





Convolution and Digital Filtering



Convolution

- **Convolution** combines two signals, **similarly to cross-correlation**
 - it's the correlation with a reversed signal

$$\text{conv}(s1, s2)[t1] = \sum_{t=0}^{N2-1} s1[t1-t]s2[t]$$

$N2$ is the length of $s2$, $s1[i] = 0$ assumed where $i < 0$ or $i \geq N$

- Often **written as $s1 * s2$**



Properties

- Convolution is **commutative**:
 $s1 * s2 = s2 * s1$
- Convolution is also **associative**:
 $(x * y) * z = x * (y * z)$
- The **length** of $s1 * s2$ is **$N1 + N2 - 1$**



Convolution Example

$$s1 = [1, 0, 2, 3, 0, 1] \quad s2 = [2, 0, 1]$$

$$1, 0, 2, 3, 0, 1$$

1 0 2	2	$(0 \cdot 1 + 0 \cdot 0 + 1 \cdot 2)$
1 0 2	0	$(0 \cdot 1 + 1 \cdot 0 + 0 \cdot 2)$
1 0 2	5	$(1 \cdot 1 + 0 \cdot 0 + 2 \cdot 2)$
1 0 2	6	$(0 \cdot 1 + 2 \cdot 0 + 3 \cdot 2)$
1 0 2	2	$(2 \cdot 1 + 3 \cdot 0 + 0 \cdot 2)$
1 0 2	5	$(3 \cdot 1 + 0 \cdot 0 + 1 \cdot 2)$
1 0 2	0	$(0 \cdot 1 + 1 \cdot 0 + 0 \cdot 2)$
1 0 2	1	$(1 \cdot 1 + 0 \cdot 0 + 0 \cdot 2)$

$$s1 * s2 = [2, 0, 5, 6, 2, 5, 0, 1]$$



Convolution Theorem

- The most **important** property of the convolution is given by the **convolution theorem**:

*A **convolution** in the **time domain***

*is **equivalent** to*

*a **multiplication** in the **frequency domain**:*

$$x * y \leftrightarrow X \cdot Y$$

$$\text{meaning: } FT(\text{conv}(x, y)) = FT(x) \cdot FT(y)$$



Digital Filters

- Sound **spectra** are **changed** by **filters**
- **STFT** manipulation and resynthesis - a form of **filtering in the frequency domain**
- **Most filtering** happens in the **time domain** by **convolution**



Linear Filters

- Linear filters **sum scaled and delayed copies** of the signal to itself (**convolution** with the **scaling factors**)
- **2 types**, depending on where they take the signal from
 - **Finite Impulse Response (FIR)** filters
(use input signal)
 - **Infinite Impulse Response (IIR)** filters
(use input & output signal)



The Order of Filters

- An **FIR** filter ***f*** of **order *k*** has this **structure**
$$f(x[n]) = b_0 x[n] + b_1 x[n-1] + b_2 x[n-2] + \dots + b_k x[n-k]$$
- An **IIR** filter ***g*** of **order *k*** has this **recursive** structure
$$g(x[n]) = + b_0 x[n] + b_1 x[n-1] + b_2 x[n-2] + \dots + b_k x[n-k] \\ - a_1 g(x[n-1]) - a_2 g(x[n-2]) - \dots - a_k g(x[n-k])$$
- or as a **difference equation**
$$y[n] = - \sum_{i=1}^k a_i y[n-i] + \sum_{i=0}^k b_i x[n-i]$$
- **a_n** and **b_n** are called **filter coefficients**

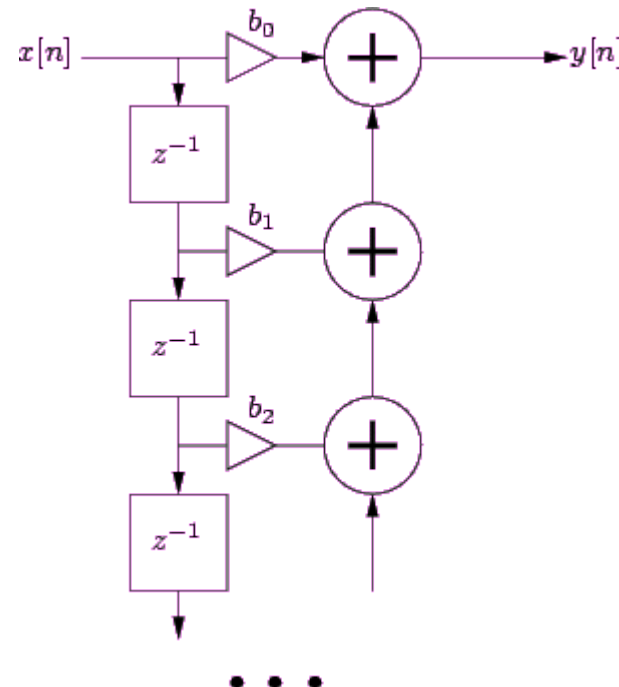


An FIR Filter

- An **FIR** filter f of order k has this **structure**

$$f(x[n]) = b_0 x[n] + b_1 x[n-1] + b_2 x[n-2] + \dots + b_k x[n-k]$$
 with **coefficients** $b = [b_0, b_1, b_2, \dots, b_k]$

- **Graphically:**



z^{-1} : delay by 1 sample



How does an FIR Filter work ?

- $f(x[n]) = b_0 x[n] + b_1 x[n-1] + b_2 x[n-2] + \dots + b_k x[n-k]$
with **coefficients** $\mathbf{b} = [b_0, b_1, b_2, \dots, b_k]$

- **FIR** is just the **convolution** $\mathbf{x} * \mathbf{b}$

$$f(x[n]) = \sum_{i=0}^k x[n-i] b[i]$$

- a small **example**:

$$\begin{array}{rcl} \mathbf{x} & = & [2, 3] \quad \mathbf{b} = [1, 2] \\ & & 2, 3 \\ & & 2 \quad 1 \quad \quad \quad 2 \\ & & 2 \quad 1 \quad \quad \quad 7 \\ & & 2 \quad 1 \quad \quad \quad 6 \\ \mathbf{f(x)} & = & [2, 7, 6] \end{array}$$



Impulse Response

- Impulse Response is **system output** in response to a **unit impulse** $[1,0,0,\dots]$.
- It **completely describes** the **behaviour** of a **linear filter** (or linear **system**) because
 - **every signal** can be described as a **sum of differently scaled unit impulses** (one per sample)
 - the **system's signal response** is then a **sum** of the differently **scaled impulse responses**
- Remember: a **linear system** satisfies the **superposition principle**
 $f(ax[n]+by[n]) = a f(x[n]) + b f(y[n])$



The Impulse Response of an FIR Filter

- The **impulse response** of a **FIR** filter

$$f(x[n]) = b_0 x[n] + b_1 x[n-1] + b_2 x[n-2] + \dots + b_k x[n-k]$$

is $h = [b_0, b_1, b_2, \dots, b_k]$, i.e. the **coefficients**

- An **example**:

$$s1 = [1, 0, 0] \quad h = [2, 0, 1]$$

$$1, 0, 0$$

$$\begin{array}{ccc} 1 & 0 & 2 \end{array} \quad \begin{array}{c} 2 \end{array} \quad f(x[0])$$

$$\begin{array}{ccc} 1 & 0 & 2 \end{array} \quad \begin{array}{c} 0 \end{array} \quad f(x[1])$$

$$\begin{array}{ccc} 1 & 0 & 2 \end{array} \quad \begin{array}{c} 1 \end{array} \quad f(x[2])$$

$$f(s1) = [2, 0, 1]$$



The Frequency Response of an FIR Filter

- The **frequency response** of an FIR filter describes its **effect on different frequency components** of a signal.
- From the **convolution theorem** we know, that the **convolution** in the **time** domain leads to a **multiplication** in the **frequency** domain: $x * h \rightsquigarrow X \cdot H$
- **H** is called the (complex) **frequency response** of the filter
 - `np.abs(H)` gives the **amplitude response**
 - `np.angle(H)` gives the **phase response**



FIR Examples

- **$b = [1, 0, 0, 0]$** does **not change** the signal, its frequency response is $[1, 1, 1, 1]$
- **$b = [0, 1, 0, 0]$** **delays** by one sample, its frequency response is $[1, -i, -1, i]$ (**magnitude stays**, but **phase changes**)
- **$b = [1, 1, 1, 1]$** (**averaging filter**) lets only the low frequencies pass, its frequency response is $[1, 0, 0, 0]$
- **Frequency resolution** depends on the **length** of the **filter**
- We can **calculate** filter **coefficients** for a given frequency response H by using `scipy.fft.irfft(H)`



Python Functions for Filters

For simple FIR filtering you can use convolution

```
y = scipy.signal.convolve(b, x)
```

The filter function lets you add IIR and truncates the output

```
y = scipy.signal.lfilter(b, a, x)
```

You can get the frequency response like this:

```
scipy.signal.freqz(b) or scipy.signal.freqz(b, a)
```

And the circular convolution treats the signals as periodic

```
y = dsp_ap.operations.circ_convolve(b, x)
```

(for an exact match between freq and time domain filtering)



Take-Home Messages

- **Short Term Fourier Transform (STFT)** for non-stationary (i.e. time varying) signals (spectrograms)
- Fourier Transform (FT) assumes **periodicity**,
 - artefacts (**spectral leakage**) for non-periodic signals
 - can be improved (but not avoided) with **Hann window**
- **Convolution** filtering (**FIR**) modifies frequency mix
- **Convolution theorem** → TD conv equiv to FT mult
- **Recurrent** filtering (**IIR**) is more **efficient**, **hard to control**
- **Impulse response** defines linear filter fully



READING

<http://www.dspguide.com/> Chapter 6 and 14.



NEXT WEEK:

Audio Programming

More Filtering