

Module IN3031 / INM378 Digital Signal Processing and Audio Programming

Tillman Weyde


t.e.weide@city.ac.uk

Acoustics Basics

- Frequency unit **Hertz** means **1/s**
- Wavelength** (λ or λ), frequency (f), **period** (p), **speed** (c) of sound ($\sim 340\text{m/s}$)
 - Example: $f = 680\text{ Hz}$, $p = 1 / 680\text{Hz} \sim 0.0015\text{ s}$,
 $\lambda = c/f = 340\text{ (m/s)} / 680\text{Hz} = .5\text{ ms/s} = .5\text{m}$
 - What is the frequency of a sound with period 10ms
- Energy** increases as **square of Amplitude** (e.g. 5-fold amplitude means 25-fold energy)
- Harmonic** sounds have components with frequencies that are integer multiples of f_0 , the lowest frequency.

Time and Frequency Domains

- Audio signals vary over time: **Time Domain**
- The spectrum decomposes the signal into frequency components: **Frequency Domain** (sin and cos)
- The **Fourier Transformation** (and its inverse) transform between Time and Frequency Domain:



Revision

Hearing

- Amplitude** - Loudness; **Frequency** – Pitch; **Spectrum** - Timbre
- Decibels**: a is $x\text{ dB}$ greater than b , means $x = 10 \log_{10} (a/b)$
- Threshold of hearing (roughly): $0\text{ dB SPL} = 10^{-12}\text{ Watt/m}^2$ and the threshold of pain 130 dB SPL
- In the inner ear, the **Cochlea** with Basilar membrane transforms **sound waves into neural signals**

Spectrum and Time

- The **spectrum** as calculated by the (fast) FT is **as long as the signal** we are analysing
 - longer signal \rightarrow higher frequency resolution
- Speeding up the signal means stretching the spectrum (and inverse for $k < 1$)

$$x(kt) \leftrightarrow \frac{1}{|k|} X\left(\frac{f}{k}\right)$$


Audio Signal Processing Summary

Digitizing Audio

- Sampling**: regularly measuring air pressure (or voltage, or any other quantity)
- Sample rate**: how many measurements (per sec)
- Sample depth: resolution** (# of bits) per sample value
- Storage space** for audio recordings:
 $\text{channels} * \text{time} * \text{sample rate} * \text{sample depth}$

Special Spectra


- Spectrum of a **unit impulse** at time 0 consists of **all real 1s** (cos), with 0 imaginary part. Real/imaginary ratio varies according to impulse's position in time.
- The spectrum of a **constant signal** with amplitude 1 is N (=length of the signal) at **frequency 0**, and 0 elsewhere.



This image shows only the lower half, the upper half is the mirror image of the lower.

Special Spectra 2

- The spectrum of a **single sinusoid** (sine/cosine mixture) with frequency f is **0 in all places but f** (only exact if the signal length is a multiple of the frequency period).
- For a **single sinusoid** with (max.) amplitude 1, the spectrum at f will have magnitude 1 (length of signal, distribution real/imaginary depends on cosine/sine ratio)



Cross-Correlation

Cross-correlation is a correlation at a lag

$$xcorr(s1, s2, k) = \sum_{t=0}^{N-1-k} s1[t] s2[t+k]$$

```
xcorr([1,2,3],[4,3,2],1)
[ 1, 2, 3]
[ 4, 3, 2]
1*0+2*4+3*3+0*3 = 17
```

Short Term Fourier Transform

- Window width**
 - frequency vs. time resolution
- Window function** enforces periodicity
 - Hann
- Crossfading**
 - necessary since signal may be not 0 at boundaries after processing
 - functions should add up to 1 in the overlapping region
 - Triangular or Hann

Sampling Theorem

- Sampling Theorem (Nyquist Theorem):**
 - When sampling with frequency f_s , only frequencies less than $f_s/2$ (**Nyquist Frequency N_y**) can be recorded accurately.
 - If this is not taken into account, **Aliasing** occurs.
- Aliasing:** a signal component with a frequency greater than N_y is reproduced as a lower frequency:

$$f_a = |((f + N_y) \bmod f_s) - N_y|$$

Other Signal Properties

Autocorrelation: Measure of self similarity, cross-correlation of a signal with itself.

$$ac(x, k) = \sum_{t=0}^{N-1-k} x[t] \cdot x[t+k]$$

Example: If signal x contains an echo at lag l autocorrelation $ac(x, l)$ will be high. If the echo is inverted ($\cdot -1$), the autocorrelation will be negative.

Sound Synthesis

- Subtractive synthesis:



Correlation

Correlation is a measure of similarity

$$corr(s1, s2) = \sum_{t=0}^{N-1} s1[t] s2[t]$$

Correlation Coefficient ρ

$$\rho = \frac{corr(x, y)}{\sqrt{\sum x[n]^2 \cdot \sum y[n]^2}}$$

Positive correlation coefficient of two signals means correlated signals (good mono compatibility of stereo signals), 0 means uncorrelated, negative correlation means cancellations in case of mono mixdown.


Filtering

- Linear Filters** add up samples with different coefficients depending on their 'age' to a new sample value
 - Finite Impulse Response **FIR** filtering applies only to input
 - Infinite Impulse Response **IIR** filtering also uses output samples (feedback loop)
- FIR filtering is *Convolution*
 - Convolution Theorem** allows calculation and design of frequency response and filter design:

$$x * y \leftrightarrow X \cdot Y$$

Sound Synthesis

- Waveforms:
 - Sine**
 - Sawtooth**
 - Square**
 - Pulse Width Modulation**
 - Narrower pulse has more harmonics (high freq)
- Need to compensate DC component



Envelop Generators

- ADSR generator.




Image Processing Summary (1)

- What is digital image processing?
- Topics in digital image processing
- Digital image acquisition
- Digital image representation
- Elementary image processing operations
- Colour image processing
- Noise in image processing
- Basic image processing in Matlab

Image Processing Summary (4)

- Spatial filtering
- Smoothing/sharpening filters
- 2-D convolution
- 2-D Discrete Fourier Transform (DFT)
- Frequency domain filters

Buffers and Latency in Digital Sound Systems

- Use buffers throughout for processing speed
- Buffers cause latency (delays between input and output)
 - Larger buffers
 - + less CPU usage, more stability
 - higher latency
 - Smaller buffers
 - more CPU, less stability
 - + lower latency

Image Processing Summary (2)

- **Digital image processing** concerns the transformation of an image to a digital format and its processing by computers
- **Topics:** image compression, medical imaging, image restoration, remote sensing, face detection...
- **Image acquisition elements:** energy, optical system, sensor
- The quality of a digital image is largely determined by the number of samples and discrete intensity levels used in **sampling** and **quantization**

Image Proc. Summary (5)

- **Mask:** small sub-image used for spatial filtering
- **Smoothing filters** are used for blurring and noise reduction
- **Sharpening filters** highlight fine detail or enhance detail that has been blurred.




Image Processing Summary

Image Processing Summary (3)

- An image can be represented by a 2-D array
- **Resolution:** spatial, gray levels
- **Basic transformations:** image negatives, gamma correction, addition, subtraction
- **Colour models:** RGB, CMY, HSI
- **Noise** in image processing: film grain, CCD noise...
- **Matlab Image Processing Toolbox**



Image Processing Summary (6)

- 2-D convolution: moving a mask from pixel to pixel
- Most filtering takes place in the frequency domain (reason: computational efficiency)

The diagram illustrates three 3x3 grids. The first grid shows a uniform value of 2 in all cells. The second grid shows a central cell (row 2, column 2) with a value of 1, surrounded by cells with a value of 2. The third grid shows a central cell (row 2, column 2) with a value of 12, surrounded by cells with a value of 6, with the outermost cells having a value of 2.

Image Processing Summary (7)

- 2-D Discrete Fourier Transform (DFT)
- Lowpass/highpass filters




Image Processing Summary (10)

- Implementation:
 - Compute the image histogram (n : total number of pixels, n_k : number of pixels with gray level r_k):

$$p_r(r_k) = \frac{n_k}{n} \quad k = 0, 1, 2, \dots, L-1$$

Compute the **cumulative histogram**:

$$s_k = \sum_{j=0}^k p_r(r_j)$$

New gray level values:

$$h(r_k) = \frac{s_k - s_{\min}}{(MN) - s_{\min}} (L-1)$$

Games Audio Summary

Image Processing Summary (8)

- 2-D correlation:

$$x[m, n] \circ y[m, n] = \frac{1}{RS} \sum_{r=0}^{R-1} \sum_{s=0}^{S-1} x^*[r, s] y[m+r, n+s]$$
- Used for template matching




Image Processing Summary (11)

- Histogram equalisation example:

52	55	63	67	63
63	59	55	90	90
63	59	68	90	68
63	58	68	55	63
67	52	68	59	55

Image

52	2	2	0
55	4	6	44
58	1	7	55
59	3	10	89
63	6	16	155
67	2	18	177
68	4	22	222
90	3	25	255

Pixels, histogram, cumulative histogram, new pixel values

Real Time DSP

- DSP algorithms can run in real time on general-purpose computers
- A lot of tasks can actually be quite cheap to run ($O(n)$ or $O(n \log n)$)
- Applications: DAWs, effects racks, VoIP and video software.

Image Processing Summary (9)

- Histogram equalisation: method for contrast adjustment

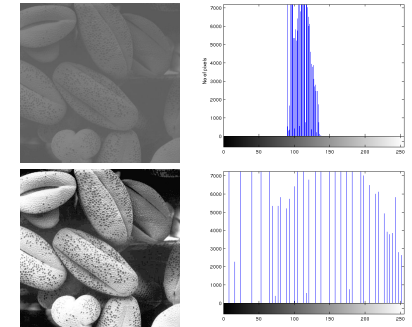



Image Processing Summary (12)

- Image restoration: reconstruct or recover an image that has been degraded

$$g(x, y) = h(x, y) * f(x, y) + \eta(x, y)$$
- Noise cannot be predicted but can be approximately described in statistical way using a probability density function (PDF)
- Minimum mean square error filtering (Wiener filtering)

The Real Time Paradigm

- Real time DSP is done by processing a **buffer** of samples.
- This means there will always be some **latency** (delay) between input and output.
- Increasing the **buffer size** increases the latency, but reduces computational cost.



Typical Sound Engine Features (e.g. FMOD)

- Support for SW and HW voices and effects
- 3D sound rendering
- Reverb and many other effects
- Dynamic voice management by audibility (Virtual Voices)
- Sound processing **graphs**
- Allow **sound designers** and **audio developers** to work independently

Doppler Effect

- Doppler Effect changes frequency for moving sources
 - $fp = f \cdot vs / (vs - vr)$
with fp: perceived frequency ,
f: frequency, vs: velocity of sound (~340m/s),
vr: velocity relative to the listener (positive = approaching)
 - Example: car moves with 68 m/s producing a 300 Hz sound. fp for a stationary listener in front of the car is
 $300 \text{ Hz} \cdot 340 / (340 - 68) = 300 \text{ Hz} \cdot 5/4 = 375 \text{ Hz}$

Programming 3D Sounds with FMOD

- Set up the FMOD systems 3D settings:
`system->set3DSettings(doppler, distFactor, distRolloff);`
- Where:
 - doppler scales the intensity of the doppler effect
 - distFactor determines the length of an FMOD unit (1 means 1m – in the OpenGL template, use 10)
 - distRolloff scales the distance roll-off (1 is like real world)

Directional Hearing and Localisation

- Directional hearing is based mostly on binaural hearing
 - Interaural Intensity Differences (IID)
 - Interaural Time Differences (ITD)
 - IID and ITD vary over frequencies.
- IID and ITD give only information on left-right
 - front-back and high-low are detected through head-related transfer functions (head shape, pinna)
- Room reflection
 - absorbing and reflecting objects give clues for source location

Programming (2D) Sounds with FMOD

- Load a sound:

```
FMOD::Sound *sound;  
result = system->createSound(filename,  
    FMOD_LOOP_OFF, 0, &sound);  
FmodErrorCheck(result);
```
- Create a Channel object and play the sound

```
FMOD::Channel channel = null;  
result = system->playSound(sound, NULL, false,  
    &channel);
```


`channel` now has the channel where sound is played

```
// set the Volume  
result = channel->setVolume(0.8f);
```

Custom FMOD DSPs

DSP inserts

- For whole system (all channels):
`system->addDSP()`
- For specific channel
`channel->addDSP()`



Intensity and Distance


- Intensity is power (energy per time) per area, measured in Watts/Meter²
- 0 dB (decibels) Sound Pressure Level defined as 10⁻¹² W/m² (~threshold of hearing)
- dBs are on a logarithmic scale
- Intensity decreases as the square of distance
- Power increases as the square of amplitude
- Example: 60dB SPL at 1m means 40dB SPL at 10m
(10-fold distance -> 100-fold decrease -> -10² = -20 = -20dB)

3D Modelling in FMOD

- FMOD has its own 3D coordinate system and model
- Need to coordinate
 - listener positions and velocities
 - sound source positions and velocities
 - objects in 3D that occlude or obstruct sound


Circular Buffer

- Avoid high buffer maintenance costs
- Address the buffer for writing:
`pos % bufferLength`
- Use index delay offset
- `(pos - delay) % bufferLength` points to read position



Occlusion

- Obstruction: obstacle blocks direct path between sound source and listener.
- Occlusion: occluder blocks all paths from the source to the listener.



Loops and Meter


- The metrical structure is normally maintained during loop playback.
- Common loop sizes are 4, 8 or 16 bars (although sometimes musical structures have different values, e.g. 'Eleanor Rigby' by the Beatles has a 5 bar structure)

MIDI vs Audio in Loops

- Audio
 - can have superior quality (e.g. recorded human performance)
 - costly and lossy change of pitch/tempo
 - changing individual notes hardly possible
 - careful planning needed, good for final production
- MIDI
 - very flexible (easy to change tempo, pitch, notes)
 - can experiment, good for developing a soundtrack

Occlusion

- Material, size and shape of the object determine how much sound (if any) reaches the listener
- Effects: drop in amplitude, loss of high frequency components (low pass filter)
- Reflected sound is still audible (with reverb)



Loops and Harmony

- Harmony describes the sounding of several (pitched) notes together
- In harmonic contexts, some notes sound consonant, others sound dissonant/inappropriate.
- Layered music loops need common harmonic structure (not true for sound loops)
- Each layer in the same harmonic pattern ensures they are musically 'compatible'

Time Series Summary

FMOD for Game Music

- Supports
 - Loop
 - Synchronisation based on beats and bars
 - Conditional transitions and repetitions
- Integrated with FMOD Event system

MIDI vs Audio

- MIDI representation
 - used mostly in music production
 - used to be applied in Games directly
- MIDI is symbolic representation
 - Advantages:
 - independent tempo and pitch
 - easy to modify for musicians
 - low data volume
 - Disadvantages
 - sound quality (depends on used sound library)

Prediction by Smoothing with Moving Average

- Moving average of span k smooths the data

$$\hat{y}_t = (y_t + y_{t-1} + \dots + y_{t-k+1}) / k$$

(\hat{y} is the prediction)

- A low pass FIR filter with coefficients $1/k, 1/k, \dots, 1/k$
- In Matlab: `filter([.25, .25, .25, .25], [1], Y)`

Exponentially Weighted Moving Average

- Moving average of infinite span smoothes the data

$$\tilde{y}_i = w y_i + (1-w) \tilde{y}_{i-1}$$
- A low-pass IIR filter with coefficients w and $(1-w)$
 In Matlab: `filter([.25],[1, -(1-.25)],Y)`
- Assumption:
 - Recent values are more important than older ones

Residual Autocorrelation

- After linear modelling and seasonal adjustment we can study the autocorrelation of the residuals $e_t = y_t - \hat{y}_t$

Correlations			
	Residual Log Passengers	Lag 1 Residuals	Lag 2 Residuals
Residual Log Passengers	1.0000	0.7896	0.6722
Lag 1 Residuals	0.7896	1.0000	0.7832
Lag 2 Residuals	0.6722	0.7832	1.0000

- With linear regression we can improve the prediction based on residuals

$$\hat{e}_t = -0.000153 + 0.7918985e_{t-1}$$

- This is a form of a generalised (weighted) moving average

$$y_i = \hat{y}_i + e_i + \sum_{l=1}^q \theta_l e_{i-l}$$

Maximum Likelihood and Regularisation

- Common approach: Linear models with least squares optimisation
- Maximise the likelihood of the data given the prediction (assuming normal distribution)
- Regularisation helps avoid overfitting (especially with small datasets)
 - Most popular: keep size of parameters low using a 'penalty term': sum of squares or absolutes of the parameters (*ridge* or *lasso*)
 - Add penalty term to errors and calculate gradient to optimise (use packaged solutions)

Linear Trend Estimation

- Linear regression: find a straight line to fit the data

$$\hat{y}_i = a_0 + a_1 i$$
- Determine a_0 and a_1 to minimise the sum or squares error


$$sse = \sum_i (\hat{y}_i - y_i)^2$$
- Solve the system of equations
 In Matlab: `coeff = polyfit(t,y,1)`

Harmonic Modelling

- Fourier modelling for residuals or seasonality
 - keep only strong components (assumption: weaker components contain noise)
 - Advantage: more efficient than autocorrelation, calculation with FFT
 - Disadvantage: need to know cycle length, less robust than autocorrelation

Cross-Validation for Regularisation

- Optimise regularisation and other parameters:
 - Divide the data into k equally sized subsets ('folds')
 - Adapt the model to $k-1$ joint subsets, test on the remaining subset, and iterate through all folds
 - Test a grid of regularisation values (or other parameters) and choose the one with best results on test sets



Seasonal Average Method

- Seasonal averages = seasonal values total / # of years
- General average = seasonal averages total / # of seasons
- Multiplicative modelling:
 Seasonal index = seasonal average / general average
- Additive modelling:
 Seasonal offset = seasonal average – general average

Modelling Caveats

- Overfitting:
 - to many parameters → model learns noise in the data but not the trend
 - need to test on data not used in building the model
 - cross-validation can when data is scarce
- Predictions get less reliable further away from sample data

More Information

- Mock paper on Moodle for your reference
- If you have any questions, please get in touch