

INM431 Machine Learning

Artur S. d'Avila Garcez

a.garcez@city.ac.uk

<http://www.staff.city.ac.uk/~aag/>

@AvilaGarcez

Based on G. Hinton's slides and C. Bishop's book

Content

Gaussian distribution

Hypothesis evaluation (in the context of the coursework task)

The Bayesian framework

Expectation and co-variance

The average value of function $f(x)$ under a probability distribution $p(x)$ is called the expectation of $f(x)$, denoted:

$$\mathbb{E}[f] = \sum_x p(x)f(x)$$

Or in the case of continuous variables:

$$\mathbb{E}[f] = \int p(x)f(x) dx$$

Given a finite number N of points drawn from the probability distribution or density then:

$$\mathbb{E}[f] \simeq \frac{1}{N} \sum_{n=1}^N f(x_n)$$

© Artur Garcez

Co-variance

$\text{cov}[x,y]$ express the extent to which random variables x and y vary together; it is a measure of linear dependence

If x and y are independent then $\text{cov}[x,y] = 0$, but the converse is not true, e.g. $y = x^2$

In the case of two vectors \mathbf{x} and \mathbf{y} , the covariance is a matrix

$$\begin{aligned} \text{cov}[x, y] &= \mathbb{E}_{x,y} [\{x - \mathbb{E}[x]\} \{y - \mathbb{E}[y]\}] \\ &= \mathbb{E}_{x,y} [xy] - \mathbb{E}[x]\mathbb{E}[y] \end{aligned}$$

$$\begin{aligned} \text{cov}[\mathbf{x}, \mathbf{y}] &= \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\{\mathbf{x} - \mathbb{E}[\mathbf{x}]\} \{\mathbf{y}^T - \mathbb{E}[\mathbf{y}^T]\}] \\ &= \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\mathbf{x}\mathbf{y}^T] - \mathbb{E}[\mathbf{x}]\mathbb{E}[\mathbf{y}^T] \end{aligned}$$

© Artur Garcez

Correlation = Normalised co-variance

The magnitude of the covariance is not necessarily easy to interpret...

Correlation:

$$\rho_{XY} = E[(X - E[X])(Y - E[Y])]/(\sigma_X \sigma_Y)$$


The covariance of a variable with itself is called its variance

The correlation of a variable with itself is always 1

© Artur Garcez

The Gaussian Distribution

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}$$



$$\mathcal{N}(x|\mu, \sigma^2) > 0$$

$$\int_{-\infty}^{\infty} \mathcal{N}(x|\mu, \sigma^2) dx = 1$$

Gaussian Mean and Variance


$$\mathbb{E}[x] = \int_{-\infty}^{\infty} \mathcal{N}(x|\mu, \sigma^2) x \, dx = \mu$$

$$\mathbb{E}[x^2] = \int_{-\infty}^{\infty} \mathcal{N}(x|\mu, \sigma^2) x^2 \, dx = \mu^2 + \sigma^2$$

$$\text{var}[x] = \mathbb{E}[x^2] - \mathbb{E}[x]^2 = \sigma^2$$

Precision parameter $\beta = 1/\text{var}[x]$



Gaussian Parameter Estimation



$$p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \mathcal{N}(x_n|\mu, \sigma^2)$$

So far...

Polynomial or Gaussian parameter estimation (choice is a **bias**)



$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

$$p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \mathcal{N}(x_n|\mu, \sigma^2)$$

© Artur Garcez

Coursework task

Choose a dataset

Classification or regression task?

Regression: predict UK Oil and Gas share price given various other share prices

Classification: predict "buy" or "sell" UK Oil and Gas shares


© Artur Garcez

Classification or regression task?

Example data:

X	Y
1.00	1.00
2.00	2.00
3.00	1.30
4.00	3.75
5.00	2.25

Generalization:
 $X = 6, Y = ?$



Regression: curve fitting, logistic regression, etc.
Classification: decision trees, naïve Bayes, etc.

© Artur Garcez

Supervised or Unsupervised Learning?

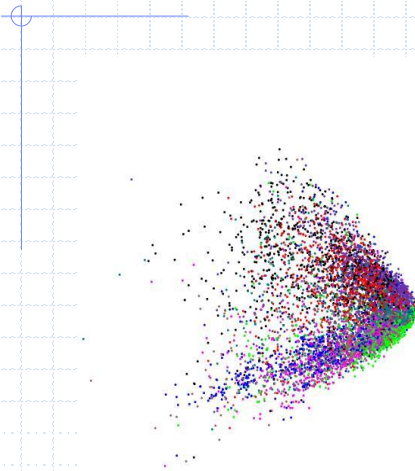
Supervised: presence of target value (or label) allows calculation of error (or loss)

Obtaining good quality labels may not be an easy task

Unsupervised: no target value known (includes reinforcement learning); learn a joint distribution given the data instead

© Artur Garcez

Example of unsupervised learning: Displaying the structure of a set of documents using Latent Semantic Analysis (a form of PCA)



Each document is converted to a vector of word counts. This vector is then mapped to two coordinates and displayed as a colored dot. The colors represent the hand-labeled classes.

When the documents are laid out in 2-D, the classes are not used. So we can judge how good the algorithm is by seeing if the classes are separated.

Semi-supervised learning

Data pre-processing (e.g. dimensionality reduction, which is unsupervised) followed by supervised learning (i.e. in the presence of labels, either a regression or classification task)

(Coursework) Include a description of the data (see example poster):
continuous/discrete (binning), how many data points, dimensions, missing values, basic statistics (correlation, skewness), data normalization (re-scaling)...


Data normalization

Rescaling: brings values into [0,1]
Scaling to unit length: dividing by magnitude
Standard score: assumes a normal distribution

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

$$x' = \frac{x}{\|x\|}$$

$$x' = \frac{x - \mu}{\sigma}$$



© Artur Garcez

Choose two methods for comparison

Popular methods:

Logistic Regression, Naïve Bayes,
Mixture of Gaussians, Hidden Markov Models,
Decision Trees, Random Forests
K-nearest neighbours (unsupervised)
Principal Component Analysis (unsupervised)

A note on sequence learning: time-series data can be rewarding to analyse but tricky to compare fairly with *sliding window* method

A note on combinations (e.g. PCA + regression) and variations (e.g. gradient boosting as variation of RFs): very good but only after the basics is done!

© Artur Garcez

Evaluation Methodology for Comparison: Generalisation

- ⊕ We expect g (our approximation of f) to produce reasonable results for examples not seen during training, i.e. we expect the learning algorithm to *generalise* to unseen cases.
- ⊕ We try to minimise a training set error (or maximise log likelihoods on the training data) as a proxy for minimising the model's generalisation error (assuming that training and test data come from the same distribution)

© Artur Garcez


Estimating Generalisation Error

- ⊕ Partition the set of examples into a training set and a validation/test set.
- ⊕ The validation/test set is not seen by the model during training.
- ⊕ The validation/test set error is an estimate of the generalisation error.
- ⊕ If the above process is repeated for different validation/test sets then your confidence on the estimate will improve.
- ⊕ Cross-validation or bootstrapping offer a systematic way of repeating this process.

© Artur Garcez

n-fold Cross-Validation

- ⊕ Divide the set of examples E into n subsets E_1, E_2, \dots, E_n
- ⊕ For each E_i ($1 \leq i \leq n$), train a model with $E - E_i$ and test it (do not change parameters) with E_i
- ⊕ Calculate the average validation/test set error



Note: Leaving m out = $(|E| / m)$ -fold cross validation.

© Artur Garcez

Bootstrapping

- ⊕ Create k (pseudo) sets of examples E_1, E_2, \dots, E_k by randomly selecting $|E|$ elements (with replacement) from E
- ⊕ For each E_i ($1 \leq i \leq k$), train a model with $E - E_i$ and test it with E_i
- ⊕ Calculate the average **test set error**; notice that the term is being used loosely, it is more precisely an averaged validation a.k.a. dev set (for development set) error

© Artur Garcez


Model Selection

- ◆ Cross-validation can also be used for model selection...
- ◆ Split examples into training and test set (e.g. 70% vs. 30% or 90% vs. 10%)
- ◆ Apply cross-validation on training set
- ◆ Use averaged validation errors to choose each model's set of hyper-parameters
- ◆ Calculate test set error of chosen models for a comparison of the ML methods

© Artur Garcez

Stopping criteria

- ◆ Training set error below 10%
- ◆ A fixed amount of time has passed
- ◆ No improvement over time
- ◆ **Early stopping** or a combination of the above



© Artur Garcez

Example

Model selection vs ensembles and grid search

Polynomial curve fitting example: parameters \mathbf{w} (the coefficients of the polynomial) and hyper-parameter M (the degree of polynomial)

Model selection: estimate \mathbf{w} for a range of values for M (e.g. $M=1,2,\dots,20$) and select the model (value of M) with the highest accuracy

Ensemble method: average the result of a number of models (i.e. combining multiple values of M)

© Artur Garcez

Example (cont.)

Grid search: if using regularization, there is also the regularization parameter λ

$M \backslash \lambda$	0.1	0.2	0.3	0.4
1				
2				
3				

ML paradox: ensemble methods tend to achieve higher accuracy than their constituent (simpler) models

© Artur Garcez

Evaluation metrics

Classification accuracy or mean-squared error
(RMSE, NMSR, etc.) $Acc = 1 - NMSE$

Average training time (given a machine)

Confusion matrix: true positives (TP), false positives (FP), true negatives (TN), false negatives (FN)

F measure and ROC curves...

n=165	Predicted:		
	NO	YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	

Other common terminology (not used by the ML expert):
sensitivity/specificity
precision/recall

© Artur Garcez


F measure and ROC curves

F1 score = $2TP/(2TP + FP + FN)$
(better than accuracy e.g. for unbalanced datasets; ranges from 0 to 1 (best), used a lot in image analysis)

Receiver Operating Characteristic: false positive rate (FP/N) vs. true positive rate (TP/P)

To compare multiple confusion matrices (binary classifiers)

AUC (area under the ROC curve) = probability that a classifier will rank a randomly chosen positive instance higher than a randomly chosen negative one



© Artur Garcez

Are you a Bayesian or Frequentist?

If we only tossed a coin once and got *heads*...

Is $p(\text{heads})=0.5$ or $p(\text{heads})=1$?

Is it reasonable to give a single answer?

- If we don't have much data, we are unsure about p .
- Our computations will work much better if we take this uncertainty into account.

The Bayesian framework

It assumes that a prior distribution always exists.


The prior may be very vague

When we see some data, we combine our prior distribution with a likelihood function to get a posterior distribution (and keep applying Bayes' theorem iteratively)


The Bayesian framework (cont.)

- ◆ Start with a prior distribution over p . In this case we used a uniform distribution

probability density




- ◆ Multiply the prior probability of each parameter value by the probability of observing *heads* given that value



- ◆ Then scale up all of the probability densities so that their integral comes to 1. This gives the posterior distribution.


probability density



Let's do it again:

- ◆ Start with a prior distribution over p .
- ◆ Suppose you get **tails** now.


probability density



- ◆ Multiply the prior of each parameter value by the probability of observing **tails** given that value.




- ◆ Then re-normalize to get the posterior distribution.




Let's do it 100 times:

- ◆ After 53 heads and 47 tails we get a very sensible posterior distribution that has its peak at 0.53 (having assumed a uniform prior)



Bayesian approach

Prior x Likelihood *is proportional to* Posterior



One step of sequential Bayesian inference: the prior is given by a Beta distribution with parameters $a = 2$, $b = 2$ in this case, and the likelihood function corresponds to a single observation of *heads*. The posterior is a beta distribution with parameters $a = 3$, $b = 2$. The parameters a and b of the Beta distribution are called *hyper-parameters* because they control the distribution of the parameter μ .

© Artur Garcez

Density estimation

We would like to model the probability distribution $p(x)$ of a random variable x given a finite set $D=\{x_1, \dots, x_N\}$ of observations (data points).

The problem of density estimation is fundamentally ill-posed because there are infinitely many probability distributions that could have given rise to the finite data set observed...

© Artur Garcez

A sampling assumption

- ◆ Assume that the training examples are drawn **i**ndependently from the set of all possible examples.
- ◆ Assume that each time a training example is drawn, it comes from an **i**dentical **d**istribution (**i.i.d**)
- ◆ Assume that the test examples are drawn in exactly the same way: i.i.d. and from the same distribution as the training data.
- ◆ These assumptions make it very unlikely that a strong regularity in the training data will be absent in the test data.

Some Loss Functions

- ◆ Squared difference between actual and target real-valued outputs (squared error).
- ◆ Number of classification errors:
 - Problematic for optimization because the derivative is not smooth.
- ◆ Negative log probability assigned to the correct answer:
 - This is usually the right function to use!
 - In some cases it is the same as squared error (e.g. regression with Gaussian output noise); in other cases it is very different...
- ◆ Cross-entropy error: use it with multi-class discrete classification tasks, i.e. one-hot target vectors

Loss function example

ML system gets the first two examples right but not the third example

ML output	Target
0.1 0.3 0.6	0 0 1
0.2 0.6 0.2	0 1 0
0.3 0.4 0.3	1 0 0

Classification error = 1/3

Squared error for first example = $(0.1 - 0)^2 + (0.3 - 0)^2 + (0.6 - 1)^2 = 0.26$

MSE = $(0.26 + 0.24 + 0.74)/3$

Least squares = $1/2 (0.26 + 0.24 + 0.74)$

Notice that the error values above are different for first and second examples, despite the same answer 0.6

Cross-entropy = $-\sum_{i=0}^n \ln(o_i) * t_i = -(\ln(0.1)*0 + \ln(0.3)*0 +$

$\ln(0.6)*1) = 0.51$ for first example. Also, 0.51 for second example!

Mean cross entropy = 0.74

Maximizing Log Likelihood

Suppose that the probability of a coin landing heads ($p(x=h)=\mu$) is not necessarily the same as that of it landing tails ($p(x=t)=1-\mu$)

We can construct a likelihood function on the assumption that the observations are drawn independently (iid):

$$p(\mathcal{D}|\mu) = \prod_{n=1}^N p(x_n|\mu) = \prod_{n=1}^N \mu^{x_n} (1 - \mu)^{1-x_n}$$

We can estimate a value for μ by maximizing the log of the likelihood function; if we set the derivative of $\log p(\mathcal{D}|\mu)$ with respect to μ equal to zero, we obtain the maximum likelihood estimator, where m is the number of *heads*:

$$\mu_{ML} = \frac{1}{N} \sum_{n=1}^N x_n = \frac{m}{N}$$

© Artur Garcez


Choosing a prior

Without a prior, suppose we flip a coin 3 times and observe 3 *heads*. Thus, $\mu_{ML}=1$

If we choose a prior to be proportional to powers of μ and $(1 - \mu)$ then the posterior distribution, which is proportional to the product of the prior by the likelihood function, will have the same functional form as the prior.

© Artur Garcez

Bayes Theorem (revisited)



Why do we maximize sums of log probabilities?

- Assuming that the errors on different training examples c are independent, we want to maximize the **product** of the probabilities of the training examples...

$$p(D | W) = \prod_c p(d_c | W)$$

- Because the log function is **monotonic**, it does not change where the maxima are. So we can maximize **sums** of log probabilities:

$$\log p(D | W) = \sum_c \log p(d_c | W)$$

- Recall $\log(xy) = \log(x) + \log(y)$
 $\log(x/y) = \log(x) - \log(y)$

Machine Learning trick 1:

- ◆ To avoid computing the posterior probabilities of all parameters W ...
 - Start with a random W and then adjust it in the direction that improves $p(W|D)$
- ◆ It is easier to work in the log domain. To maximize $p(W|D)$, minimize the *Cost* :

$$p(W | D) = p(W)p(D | W) / p(D)$$

$$Cost = -\log p(W | D) = -\log p(W) - \log p(D | W) + \log p(D)$$

Using log probabilities helps keep the values low by replacing the product of the probabilities by the sum of the log probabilities

Machine Learning trick 2:

- ◆ Completely ignore the prior over W :
 - This is equivalent to giving all possible W the same prior probability density
- ◆ Then all we need to do is maximize:

$$\log p(D | W) = \sum_c \log p(D_c | W)$$

- ◆ This is called **maximum likelihood** learning. It is widely used for fitting models in statistics...

So, Maximum Likelihood...

We want to determine the values of μ and σ that maximize the log likelihood:

$$\ln p(\mathbf{x}|\mu, \sigma^2) = -\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi)$$

Maximizing w.r.t. μ :

$$\mu_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N x_n$$


Maximizing w.r.t. σ^2 :

$$\sigma_{\text{ML}}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \mu_{\text{ML}})^2$$

Properties of μ_{ML} and σ_{ML}^2

$$\mathbb{E}[\mu_{\text{ML}}] = \mu$$

$$\mathbb{E}[\sigma_{\text{ML}}^2] = \left(\frac{N-1}{N}\right) \sigma^2$$



On average, maximum likelihood will obtain the correct mean, but will underestimate the variance by a factor $N-1/N$.

Trading off the goodness of fit against the complexity of the model

- ◆ You can only expect a model to generalize well if it explains the data surprisingly well given the complexity of the model.
- ◆ If the model has as many degrees of freedom as the data, it can fit the data perfectly but so what?
- ◆ Learning theory seeks to measure the model complexity and how to control it to optimize generalization, c.f. Les Valiant's PAC learnability.
- ◆ The main principle is known as **Ockham's razor** (William of Ockham (c. 1287–1347) was an English friar, philosopher and theologian): Among competing hypotheses, under uncertainty, the one with the fewest assumptions should be selected.